

Project 2: Hate Speech on Instagram

HON2200

Amélie Wessel
Eman Ghazniwal
Ådne Rolstad
Marius Torsheim

Universitet i Oslo
Mai 2026
Gruppe 12

Abstract

This article examines the impact of Meta’s January 2025 revision of its hate speech policy on Instagram. The policy, renamed from *Hate Speech Policy* to *Hateful Conduct Policy*, relaxed restrictions on several categories of speech, particularly those targeting gender identity, sexual orientation, and immigration status. We approach this question from two complementary perspectives: a linguistic analysis of the terminological shift from *hate speech* to *hateful conduct*, and an empirical analysis of Meta’s Community Standards Enforcement Report data from 2020 to 2026. The linguistic analysis suggests that while the shift to *conduct* implies greater emphasis on individual intent, its practical effect has been a relaxation of enforcement. This is supported by a sharp and sustained decline in content actioned for hateful conduct following the policy change in Q4 2024. Analysis of annotated hate speech data from X further shows that the groups most frequently targeted by hate speech online correspond directly to those for whom Meta relaxed its moderation standards. We also consider the jurisdictional implications of the policy change, noting that content now permitted under Meta’s revised guidelines may constitute illegal hate speech under European law. Taken together, these findings suggest that the 2025 revision disproportionately reduces protection for groups already most exposed to hateful content online.

Contents

1	Introduction	3
2	Background	3
2.1	Semantic Definition of Hate Speech	3
2.2	Metas Definition of Hate Speech	4
2.3	EU vs. USA	4
2.4	Meta Guideline Changes	5
2.5	Semantic Definition of Hateful Conduct vs. Hate Speech	5
3	Methods	6
3.1	Analytical Approach	6
3.2	Metas Community Standards Enforcement Report	6
3.3	X Hate Speech Data	7
4	Data Analysis	7
4.1	Meta Community Standards Enforcement Report	7

4.2	X Hate Speech	9
4.2.1	Davidson Dataset	9
4.2.2	HateXplain Dataset	9
5	Discussion	10
5.1	Linguistic Analysis of Hate Speech Features	10
5.2	Methodological neutrality and interpretive uncertainty	11
5.3	Criticism and Metas Response	12
5.4	Jurisdictional Consequences of Meta’s Policy Change	12
6	Conclusion	13

1 Introduction

Social media platforms have become central arenas for political discussion, identity expression, and public discourse. At the same time, they have become increasingly dangerous spaces where hate speech, harassment, and discriminatory rhetoric can spread rapidly, often with limited oversight (Christie et al., 2024).

Instagram, owned by Meta, regulates harmful speech through its Community Standards. However, these guidelines, particularly those concerning hate speech, underwent significant revisions in January 2025. The policy was renamed from *Hate Speech Policy* to *Hateful Conduct Policy*, and restrictions on certain categories of speech were loosened. These changes have sparked considerable criticism, with many arguing that the new regulations blur the line between legitimate discourse and harmful rhetoric, and that they disproportionately expose vulnerable groups to harassment.

This article examines whether Meta’s revised guidelines are sufficient to protect vulnerable groups on Instagram. We approach this question from two complementary perspectives: a linguistic analysis of the terminological shift from *hate speech* to *hateful conduct*, and an empirical analysis of Meta’s enforcement data before and after the policy change. We also consider the broader regulatory context, including the divergence between European and American approaches to hate speech, and the political pressures that appear to have shaped Meta’s decision.

2 Background

2.1 Semantic Definition of Hate Speech

Hate speech can be separated into generalized and direct forms. Directed hate speech is personal, attacking individuals, while generalized hate speech is impersonal, attacking groups as a collective or as an entity. In generalized hate speech the user often broadens the scope of the act, encompassing larger groups and escalating through more violent verbs such as kill, murder. This escalation is promoted by the depersonalization of the intended target, as it removes a directly impacted victim.

The detection of hate speech is extremely challenging. Hate speech is not defined by the presence of specific words which makes its detection more complicated than flagging vocabulary as hateful. Even in instances where slurs, words defined as having hateful value, have been used, the text may be a critique of hate speech rather than an expression of it. There is also the difficulty of the blurred lines between protected speech and hate speech (ElSherief et al., 2018).

Semantic change, the process by which the meanings of words change over time, makes using lexical detection systems difficult. Some slurs have been ‘reappropriated’ by targeted groups, who reclaim the words and use it in a positive way. This can be seen in the positive uses of the word ‘queer’ by the LGBTQ2+ community and the word ‘bitch’ by women. The meaning of words can also become more negative, pejoration. Words that were previously neutral can be used as coded language, algospeak, to mean something hateful. In this way semantic change argues for the loosening of regulation as fewer words can be regarded as outright hateful. As language is infinite and in constant change, any static system used to detect hateful uses of it could never be fully accurate or comprehensive.

2.2 Metas Definition of Hate Speech

Meta defines hate speech as “direct attacks against people rather than concepts or institutions, on the basis of what we call protected characteristics (PCs): race, ethnicity, national origin, disability, religious affiliation, caste, sexual orientation, sex, gender identity, and serious disease” (Meta Platforms, Inc., 2025b).

There is no universal definition of hate speech under international human rights law. However, Meta’s definition places stronger emphasis on direct attacks. Whereas broader international understandings, such as those used by the United Nations, tend to be more inclusive, covering discriminatory rhetoric targeting individuals or groups based on identity factors (United Nations, 2023). For the sake of clarity, this article will use Meta’s definition of hate speech, which is almost identical to the EU definition of illegal hate speech.

2.3 EU vs. USA

It is important to consider the differences between how the EU and the United States define and regulate hate speech, as these differences influence how Meta’s guidelines are implemented across jurisdictions.

The United States is highly protective of freedom of speech, as engrained in the First Amendment of the Constitution, although this law regulates infringement by the government on speech, preventing it from restricting the free speech of the people, and does not apply to the actions of private companies. Social media platforms can prevent speech acts as much or as little as they like under the First Amendment. So if the platform removes a post, it is because the post has violated their guidelines and is not necessarily a First Amendment violation. Even offensive or hateful speech is generally allowed, as long as it does not pose an immediate threat, constitute illegal activity, or incite imminent violence, depending on the private moderation of the private companies themselves. In this context, censorship of ideas is often viewed as a greater threat than the harm arising from hate speech. Social media platforms are also protected under Section 230 of the Communications Decency Act of 1996, which holds users—not platforms—responsible for the content they publish. (ScienceDirect, 2024). As a result, government involvement is limited, and platforms have broad discretion to moderate content voluntarily.

In contrast, the European Union adopts a more restrictive approach. It views freedom of expression as a right that comes with responsibilities and limitations. The EU is more willing to restrict speech if it threatens public safety, social order, or the rights of others. This approach is shaped by historical experiences, particularly the Holocaust and totalitarian regimes, where speech was used to justify and promote violence and discrimination (Car and Immenkamp, 2025).

Under the 2008 Framework Decision, hate speech that targets gender, sexuality, and disability is criminalized. Additionally, the Digital Services Act (DSA) requires platforms to remove illegal content when reported, implement reporting mechanisms, and assess systemic risks, particularly on large platforms (Car and Immenkamp, 2025). Unlike the United States, the EU holds platforms themselves accountable for the content they host.

The president of the United States has criticized the EU for potentially censoring ideas and opinions, especially in the digital space. Meanwhile, companies such as Meta, X, and YouTube define their own hate speech policies in the U.S., often influenced by political and social debates. However, when operating in the EU market, these companies must comply with European laws and regulations.

2.4 Meta Guideline Changes

In January 2025, Meta revised its Hate Speech Policy and renamed it the Hateful Conduct Policy. These changes made the guidelines more permissive compared to previous versions, particularly by relaxing restrictions related to gender identity, immigration, and sexuality.

The guidelines still operate through a three-tier system of severity, but the definitions within each tier have changed. Exclusionary language may be permitted if framed within political or social debate, including discussions on sensitive topics such as immigration and refugee policy. Some insulting or exclusionary speech is now allowed, especially in political or religious discourse. This is something that was previously restricted.

- **Tier 1** represents the most severe forms of hate, associated with the highest risk of real-world harm. This includes slurs, dehumanization, comparisons to animals or disease, and calls for violence or extreme exclusion. However, some content that previously fell under Tier 1 has been moved to lower tiers. Content categorized in Tier 1 will be removed from Meta’s platforms. For example, it is now allowed in some contexts to describe LGBTQ+ identities as “abnormal” or “mentally ill” which were previously categorized at Tier 1 are now pushed down to Tier 2. Statements denying transgender identities are also no longer categorically restricted or automatically removed.
- **Tier 2** includes more implicit forms of harmful speech, such as coded or indirect expressions of exclusion or segregation. Before 2025, this tier covered insults, stereotypes, and generalizations about protected groups, many of which were removed. After the guideline changes, the scope has widened: some insults and generalizations, particularly in political or religious contexts, are now allowed. For instance, statements like “most migrants are criminals” remain banned, but “some migrants are criminals” may now be permitted. Not all content categorized in Tier 2 will be removed from Meta’s platforms.
- **Tier 3** includes indirect or less severe harmful content and functions as a buffer zone where borderline content is often allowed, especially when framed as opinion or debate. Most content in Tier 3 will be permitted on Meta’s platforms.

A key change across the system is that it has become harder for content to qualify as Tier 1, meaning more content now falls into Tier 2 or Tier 3 and is therefore more likely to remain on the platform.

2.5 Semantic Definition of Hateful Conduct vs. Hate Speech

The shift from *speech* to *conduct* broadens the language used to define the policy. *Speech* is limited to expressions through words or sounds, whereas *conduct* encompasses any act or behaviour. This broadening places greater emphasis on the person behind the act rather than the act itself, shifting the definition from characterising the expression as hateful to characterising the actor as hateful through their actions. In turn, this increases the culpability of the individual behind the act.

This reframing also reinforces the distinction between generalised and direct hate speech. With greater emphasis placed on personal intention, direct hate speech targeting a specific individual becomes more readily identifiable as a violation, since intent is more apparent. Generalised hate speech, which targets groups rather than individuals, may be harder to prosecute under a conduct-based framework, as the personal intent behind such speech is less explicit.

In practice, however, the policy change has reduced the severity assigned to specific acts, with certain content that previously qualified as Tier 1 now falling into lower tiers. This suggests that while the terminological shift to *conduct* implies greater accountability, the practical effect of the 2025 revisions has been a relaxation of enforcement.

3 Methods

3.1 Analytical Approach

To assess the impact of Meta’s January 2025 policy change, we conduct a before-and-after analysis of Instagram’s hateful conduct enforcement data from Q1 2020 to Q1 2026, with the policy change in Q4 2024 serving as the primary point of comparison. We visualise all four metrics reported by Meta as time series plots, allowing for visual inspection of trends and discontinuities around the policy change. In addition, we compute descriptive statistics for the periods before and after the policy change to quantify differences in enforcement activity.

In the current analysis, we focus primarily on content actioned, as it provides the most direct measure of enforcement.

It should be noted that all data is self-reported by Meta, and changes in the metrics may reflect changes in Meta’s definitions and classification practices rather than genuine shifts in the prevalence of hateful conduct on the platform.

3.2 Metas Community Standards Enforcement Report

We analyse the data from Metas “Community Standards Enforcement Report”, which contains information about how much content has been removed (actioned) by Meta on different platforms. In this project we focus on Instagram. The report contains information on the type of content that has been actioned, for instance “Fake Accounts”, “Bullying and Harassment” and “Hateful Conduct”. In addition, Meta reports the number of contents that has been restored, both with and without an appeal from a user. The data is organized as quarterly figures from 2019 to 2026.

The relevant metrics reported by Meta are described below:

- **Content Actioned** refers to the number of contents that has been removed from the platform in the relevant period. This includes content that has been reported by users, as well as content found by Metas algorithms.
- **Appealed Content** refers to actioned content that has been appealed by a user. This does not mean that the content is restored.
- **Restored Content** refers to content that has been restored to the platform. Meta distinguishes between content that has been restored after appeal, and content that has been restored automatically.
- **Prevalence** considers all the views of content on the platform, and estimates the percentage of those views that was of violating content.

There are some significant limitations to the data set. First of all, the data is very coarse-grained, showing only quarterly results. Second, there is no distinction between types of content. The term “content” could refer to photos, comments, videos, posts or users. Third, Meta does not report the number of content published in the relevant period, which makes it difficult to compute the rates of false positives and false negatives. Furthermore, the prevalence metric only measures the views of content that violates Metas community guidelines. This means that Meta themselves decides what constitutes a violation, and changes to these guidelines makes it difficult to compare the prevalence over time. This raises a deeper question about the definition of hate speech, and whose definition one should use. Finally, all metrics are self-reported by Meta with no independent verification, which creates an inherent conflict of interests.

3.3 X Hate Speech Data

To better understand the linguistic characteristics of hate speech, we analyse annotated social media posts from X (previously Twitter). Ideally, such a dataset would consist of Instagram content. However, no openly available annotated datasets of this kind exist for Instagram. The data from X provides a suitable substitute, as the short-form, public posts are comparable to the captions and comments of Instagram posts.

We use two complementary datasets. The first, compiled by Davidson et al. (2017), was constructed by scraping tweets containing terms from a curated hate speech lexicon sourced from Hatebase.org. The dataset consists of 24,783 tweets, each annotated by at least three workers recruited through the CrowdFlower crowdsourcing platform. Annotators were asked to classify each tweet as *hate speech*, *offensive language*, or *neither*. The final label for each tweet was determined by majority vote. The resulting distribution is highly skewed: 77.4% of tweets are labelled offensive language, 16.8% neither, and only 5.8% hate speech. The data was collected in 2013.

The second dataset, Mathew et al. (2021), consists of 20,148 posts drawn from both Twitter ($n = 9,027$) and the far-right platform Gab ($n = 11,093$). Each post was annotated by three workers via Amazon Mechanical Turk with one of three labels: *hatespeech*, *offensive*, or *normal*. In contrast to the Davidson dataset, annotators were also asked to identify the *target community* of the post (e.g., African, Women, Homosexual, Islam) and to mark the specific tokens that motivated their labelling decision, known as *rationales*. The majority label determines the final class, yielding 6,234 hate speech posts (31.0%), 5,761 offensive (28.6%), and 8,153 normal (40.5%). The posts were collected between 2019 and 2020.

The goal of this analysis is two-fold: we aim to identify which words and phrases (features) are most strongly associated with hate speech, and to assess which groups and minorities are most frequently targeted. We rely on the assumption that findings from X are broadly applicable to Instagram, given the structural similarities between the platforms’ short-form, public-facing content.

The analysis proceeds as follows for both datasets. Texts are first preprocessed by converting to lowercase, removing URLs, mentions, hashtags, retweet markers, punctuation, digits, and common English stopwords. The cleaned texts are then converted to numerical features using TF-IDF vectorisation with both unigrams and bigrams (maximum 20,000 features). We train a logistic regression classifier with balanced class weights to account for the skewed class distribution, and evaluate it on a held-out test set using the *accuracy score*. Since offensive language constitutes the large majority of posts in both datasets, accuracy will be inflated relative to what a model needs to achieve to be useful in practice. To interpret the model, we extract the 15 most predictive features per class from the logistic regression coefficients.

4 Data Analysis

4.1 Meta Community Standards Enforcement Report

Figure 1 shows the quarterly number of content actioned for hateful conduct on Instagram from 2019 to 2025. The series exhibits two notable peaks: one in 2021Q2 and one in 2023Q2, both approaching 10^7 removals per quarter. Following Meta’s policy change in 2024Q4, marked by the red dashed line, the number of actioned content drops sharply and consistently, reaching its lowest point in 2025Q4 at approximately 1.3×10^6 . This decline should not be interpreted straightforwardly as a genuine reduction in hateful content on the platform. Since the metric

depends on Meta’s own definitions and enforcement practices, the decline may instead reflect a narrower threshold for what Meta classifies as a violation, illustrating the definitional challenge discussed in section 2.




Figure 1: Plot showing the quarterly number of content on Instagram that has been actioned (removed) by Meta for *hateful conduct*.

For comparison, fig. 2 shows the same plot of actioned content, for the Facebook platform. This platform uses the same community guidelines, and we observe a similar drop in content actioned as for Instagram, though not as apparent.




Figure 2: Plot showing the quarterly number of content on Facebook that has been actioned by Meta for *hateful conduct*.

We also wish to compare these findings with other types of content that violates Metas community guidelines. In fig. 3, we see the quarterly number of content on Instagram actioned for “Adult Nudity & Sexual Activity”. The number of actioned content is decreasing after the policy change, but the decrease is not as sharp as for “hateful conduct”.




Figure 3: Content actioned for “Adult Nudity & Sexual Activity” on Instagram.

4.2 X Hate Speech

4.2.1 Davidson Dataset

We train a logistic regression classifier, evaluating in on a held-out test set comprising 20% of the data. The logistic regression achieves an accuracy score of 86.1%.

Figure 4 shows the 15 most predictive features per class derived from the logistic regression coefficients.




Figure 4: Top 15 most predictive features per class for the Davidson dataset, based on logistic regression coefficients. Features are ranked by their contribution to each class.

4.2.2 HateXplain Dataset

On the HateXplain dataset, the logistic regression achieves an accuracy score of 64.3%. Performance is worse than on the Davidson dataset, reflecting the increased difficulty of three-way classification on a more heterogeneous dataset spanning both Twitter and Gab. Figure 5 presents the top 15 features per class. Figure 6 shows the distribution of target groups among posts labelled as hate speech in the HateXplain dataset. African Americans, Muslims and Jews constitute the three most frequently targeted groups, followed by women, Arabs and homosexuals. Notably, several of the most targeted groups — including homosexual individuals, refugees, and Arab and Muslim communities — correspond directly to the groups for which




Figure 5: Top 15 features per class for the HateXplain dataset, based on the logistic regression coefficients. Features are ranked by their contribution to each class.

Meta relaxed its hate speech policy in January 2025 Meta Platforms, Inc. (2025a). This suggests that the platform’s policy change disproportionately reduces moderation pressure on content targeting groups that are already among the most exposed to hate speech online.




Figure 6: Top 15 target groups among hate speech posts in the HateXplain dataset.

5 Discussion

5.1 Linguistic Analysis of Hate Speech Features

In our data analysis of the Davidson and HateXplain dataset in section 4.2, we were given an insight into which words were the most predictive features in language being flagged. This shows two things: which groups are the most targeted and how severely certain vocabulary is treated. The distinction between “hate speech” and “offensive language” is whether the definition of the word changes with context: whether the semantic definition is different from the pragmatic definition. Here very few words can be considered outrightly offensive in any context: only slurs can be considered semantically hateful. The higher accuracy of the Davidson dataset is likely due to a larger proportion of slurs (words that are inherently hateful).

This has direct implications for content moderation. A system that relies on lexical detection will perform well when slurs are present, but will struggle with generalised hate speech that

avoids explicit slurs in favour of coded language or dog whistles. As discussed in section 2, semantic change means that algospeak and reappropriated terms continuously challenge static moderation systems. The lower accuracy of the HateXplain classifier reflects this: spanning both Twitter and Gab, the dataset contains more ideologically coded language that is harder to detect without contextual understanding.

The target group distribution in fig. 6 further illustrates that hate speech is not evenly distributed. African Americans, women, and Muslims are disproportionately targeted, and several of these groups, including homosexual individuals, refugees, and Arab and Muslim communities, are precisely those for whom Meta relaxed its moderation standards in January 2025. This alignment raises concern, although it should be interpreted cautiously: the overlap suggests that some groups frequently targeted in annotated hate speech data may also be affected by Meta’s relaxed moderation standards.

5.2 Methodological neutrality and interpretive uncertainty

A central limitation of this project is that the analysis depends on categories that are not entirely neutral or self-evident. Hate speech is not only a linguistic phenomenon, but also a legal, political, and social category. As discussed in section 2, hate speech cannot be identified simply by detecting specific words, since the same expression may function differently depending on context, intention, target group, speaker position, and platform norms. This means that any empirical analysis of hate speech also depends on how the concept is operationalised.

This limitation is particularly important for our analysis of Meta’s Community Standards Enforcement Report. In section 3, we use “content actioned” as the main indicator of enforcement because it is the most direct measure available in Meta’s data. However, content actioned is not the same as the actual amount of hate speech on Instagram. A decline in actioned content after the policy revision may indicate that Meta is enforcing its rules less strictly, but it may also reflect changes in classification practices, reporting behaviour, automated detection systems, or the threshold for what counts as a violation. Since the data is self-reported by Meta, lacks independent verification, and does not include the total amount of content posted, we cannot calculate false positive or false negative rates. The analysis therefore shows a change in Meta’s enforcement outcomes, but it cannot by itself prove whether hateful content became more or less common on the platform.

The same caution applies to the X hate speech datasets used in section 4.2. These datasets are useful for studying linguistic patterns in hate speech, but they are not direct evidence of hate speech on Instagram. The Davidson dataset was collected from Twitter in 2013 using a hate speech lexicon, which may overrepresent explicit slurs and offensive vocabulary (Davidson et al., 2017). The HateXplain dataset includes both Twitter and Gab posts, and therefore contains platform-specific and politically coded language that may not correspond directly to Instagram use (Mathew et al., 2021). In addition, our preprocessing removes elements such as hashtags, mentions, punctuation, and other contextual signals that may be important for interpreting social media posts. The TF-IDF and logistic regression approach can identify words and phrases associated with annotated hate speech, but it cannot fully capture irony, memes, visual context, quotation, reappropriation, or political framing.

This creates a risk of interpretive bias. If offensive, provocative, or politically controversial language is treated as equivalent to hate speech, the analysis may overstate the harm of some forms of expression. Political criticism, religious disagreement, satire, and debate about immigration or gender policy may be harsh or uncomfortable without necessarily constituting hate speech. This is the concern Meta itself raises when it argues that previous moderation

systems removed too much legitimate content. From this perspective, reducing false positives can be understood as a positive aim, because over-moderation may limit political debate and remove speech that should remain visible in a democratic public sphere.

At the same time, focusing only on false positives creates the opposite risk. If the threshold for hate speech is narrowed too much, harmful content may remain online even when it contributes to exclusion, intimidation, or discrimination against protected groups. This is especially relevant because the target group analysis in fig. 6 shows that some of the groups frequently targeted in annotated hate speech data overlap with groups affected by Meta’s revised policy categories. However, this overlap should be interpreted cautiously. It suggests a possible risk that reduced enforcement may affect already targeted groups, but it does not allow us to measure that effect directly on Instagram.

A more neutral interpretation of our findings is therefore that Meta’s 2025 policy revision coincides with a measurable decline in content actioned for hateful conduct, while the meaning of this decline remains contested. It may reflect a correction of overinclusive moderation, a reduction in protection against harmful speech, or both at the same time. The central issue is not simply whether Meta allows “more speech”, but how the platform distinguishes between legitimate disagreement, offensive expression, and discriminatory or dehumanising attacks. Recognising this uncertainty strengthens the analysis, because it makes clear that our conclusions depend on the definitions, datasets, and methodological choices used in the project.

5.3 Criticism and Metas Response

Unsurprisingly, Meta has faced significant criticism for shifting from a model of “broad moderation” to one that allows “more speech”. Critics argue that these changes disproportionately affect vulnerable groups, exposing them to higher levels of harassment and discrimination. They claim that relaxing hate speech rules may normalize harmful rhetoric, as such content is no longer consistently removed.

Meta, however, argues that the changes were necessary because the previous system removed too much legitimate content. The company frames this as a “false positive” problem, where harmless speech was mistakenly taken down.

Critics, on the other hand, argue that the issue has now become a “false negative” problem, where too much harmful content remains online without consequences. This is inherently difficult to quantify, as content that falls below Meta’s revised thresholds will no longer be actioned regardless of whether it is reported.

Another justification provided by Mark Zuckerberg is the desire to enable broader political debate on Meta’s platforms (Kaplan, 2025). As a U.S.-based company, Meta’s approach reflects a strong commitment to free speech and skepticism toward censorship, even when the content in question may be offensive or harmful by broader definitions of hate speech. At the same time, the timing and framing of the policy change suggest that Meta’s approach may also reflect broader political and regulatory pressures in the United States (Associated Press, 2025).

5.4 Jurisdictional Consequences of Meta’s Policy Change

Meta’s January 2025 policy revision reflects a broader shift towards the American free speech tradition, in which restricting expression is often viewed as a greater threat than the harm arising from hate speech itself. As noted in section 2.3, this stands in contrast to the EU’s more restrictive approach, shaped by historical experiences with speech used to justify violence and discrimination (Car and Immenkamp, 2025).

This divergence has concrete legal implications. Meta’s revised guidelines permit certain content that may nonetheless constitute illegal hate speech under the EU’s 2008 Framework Decision and the Digital Services Act (Car and Immenkamp, 2025). This is particularly notable given that Meta’s previous definition of hate speech was nearly identical to the EU’s definition, as discussed in section 2.2. The 2025 revision has therefore widened the gap between what Meta allows and what European law requires.

Unlike the American Section 230 framework, which holds users rather than platforms responsible for content, the DSA places accountability directly on platforms (Car and Immenkamp, 2025). It remains an open question whether Meta will apply its revised standards uniformly across jurisdictions, or maintain stricter moderation in the EU market to comply with local law. This question is particularly relevant for Norwegian and other European users, whose legal protections against hate speech exist independently of Meta’s internal policy choices.

6 Conclusion

This article has examined Meta’s 2025 revision of its hate speech policy from two complementary perspectives: a linguistic analysis of the terminological shift from *Hate Speech Policy* to *Hateful Conduct Policy*, and an empirical analysis of Meta’s Community Standards Enforcement Report data for Instagram from 2020 to 2026.

The linguistic analysis reveals that the shift from *speech* to *conduct* carries meaningful semantic implications, placing greater emphasis on the intention of the actor rather than the nature of the expression. In practice, however, the policy change has relaxed enforcement, with a sharp and sustained decline in content actioned for hateful conduct following the revision in Q4 2024.

The analysis of hate speech data from X further contextualises these findings. The groups most frequently targeted in the HateXplain dataset, including homosexual individuals, refugees, and Muslim communities, correspond directly to those for whom Meta relaxed its moderation standards in January 2025. This convergence suggests that the policy change disproportionately reduces protection for groups already most exposed to hateful content online.

Whether the decline in content actioned reflects a genuine reduction in hate speech or merely a narrowing of what Meta classifies as a violation remains difficult to determine. The data is self-reported, lacks independent verification, and changes in enforcement metrics are confounded by changes in the underlying definitions. It is therefore too early to draw firm conclusions about whether the new guidelines are sufficient to protect vulnerable groups on Instagram.

What the data does suggest is that definitional choices have measurable consequences for enforcement. As the regulatory landscape continues to diverge between the EU and the US, and as Meta’s policies increasingly reflect broader political pressures, the question of whose definition of hate speech governs online discourse remains an open and urgent one.

References

Associated Press. Meta rolls back hate speech rules as Zuckerberg cites ‘recent elections’ as a catalyst. January 2025. URL <https://www.wsaj.com/2025/01/09/meta-rolls-back-hate-speech-rules-zuckerberg-cites-recent-elections-catalyst/>.

- Polona Car and Beatrix Immenkamp. Hate speech: Comparing the US and EU approaches. Briefing EPRS_BRI(2025)772890, European Parliamentary Research Service (EPRS), June 2025. URL [https://www.europarl.europa.eu/thinktank/en/document/EPRS_BRI\(2025\)772890](https://www.europarl.europa.eu/thinktank/en/document/EPRS_BRI(2025)772890).
- Christina Christie, Anne Gilliland, Christine Ong, Mark Hansen, Arif Amlani, and Marc Berkman. The rise of social media hate: SMASH project findings, 2024. URL <https://studyofhate.ucla.edu/smash-social-media-hate/>. Social Media and the Spread of Hate (SMASH) Project, UCLA Initiative to Study Hate.
- Thomas Davidson, Dana Warmley, Michael Macy, and Ingmar Weber. Automated hate speech detection and the problem of offensive language. In *Proceedings of the 11th International AAI Conference on Web and Social Media*, ICWSM '17, pages 512–515, 2017.
- Mai ElSherief, Vivek Kulkarni, Dana Nguyen, William Yang Wang, and Elizabeth Belding. Hate lingo: A target-based linguistic analysis of hate speech in social media. In *Proceedings of the International AAI Conference on Web and Social Media*, volume 12, June 2018. doi: 10.1609/icwsml.v12i1.15041. URL <https://ojs.aaai.org/index.php/ICWSM/article/view/15041>.
- Joel Kaplan. More speech, fewer mistakes, January 2025. URL <https://about.fb.com/news/2025/01/meta-more-speech-fewer-mistakes/>.
- Binny Mathew, Punyajoy Saha, Seid Muhie Yimam, Chris Biemann, Pawan Goyal, and Animesh Mukherjee. HateXplain: A benchmark dataset for explainable hate speech detection. In *Proceedings of the AAI Conference on Artificial Intelligence*, volume 35, pages 14867–14875, 2021.
- Meta Platforms, Inc. Community standards enforcement report, 2025a. URL <https://transparency.meta.com/reports/community-standards-enforcement/>.
- Meta Platforms, Inc. Hateful conduct, 2025b. URL <https://transparency.meta.com/policies/community-standards/hateful-conduct/>.
- ScienceDirect. Communication decency act: An overview, 2024. URL <https://www.sciencedirect.com/topics/computer-science/communication-decency-act>. ScienceDirect Topics, Computer Science.
- United Nations. What is hate speech?, 2023. URL <https://www.un.org/en/hate-speech/understanding-hate-speech/what-is-hate-speech>. UN Strategy and Plan of Action on Hate Speech.